

In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
```

In [2]:

```
df=pd.read_csv("heart_failure_clinical_records_dataset.csv")
```

In [3]:

```
df
```

Out[3]:

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_pressure	platelets	serum_creatinine	serum_sodium	sex
0	75.0	0	582	0	20	1	265000.00	1.9	130	
1	55.0	0	7861	0	38	0	263358.03	1.1	136	
2	65.0	0	146	0	20	0	162000.00	1.3	129	
3	50.0	1	111	0	20	0	210000.00	1.9	137	
4	65.0	1	160	1	20	0	327000.00	2.7	116	
...
294	62.0	0	61	1	38	1	155000.00	1.1	143	
295	55.0	0	1820	0	38	0	270000.00	1.2	139	
296	45.0	0	2060	1	60	0	742000.00	0.8	138	
297	45.0	0	2413	0	38	0	140000.00	1.4	140	
298	50.0	0	196	0	45	0	395000.00	1.6	136	

299 rows × 13 columns

In [4]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 299 entries, 0 to 298
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                   299 non-null    float64
1   anaemia                             299 non-null    int64
2   creatinine_phosphokinase            299 non-null    int64
3   diabetes                            299 non-null    int64
4   ejection_fraction                   299 non-null    int64
5   high_blood_pressure                 299 non-null    int64
6   platelets                           299 non-null    float64
7   serum_creatinine                    299 non-null    float64
8   serum_sodium                        299 non-null    int64
9   sex                                 299 non-null    int64
10  smoking                             299 non-null    int64
11  time                                299 non-null    int64
12  DEATH_EVENT                         299 non-null    int64
dtypes: float64(3), int64(10)
memory usage: 30.5 KB
```

In [5]:

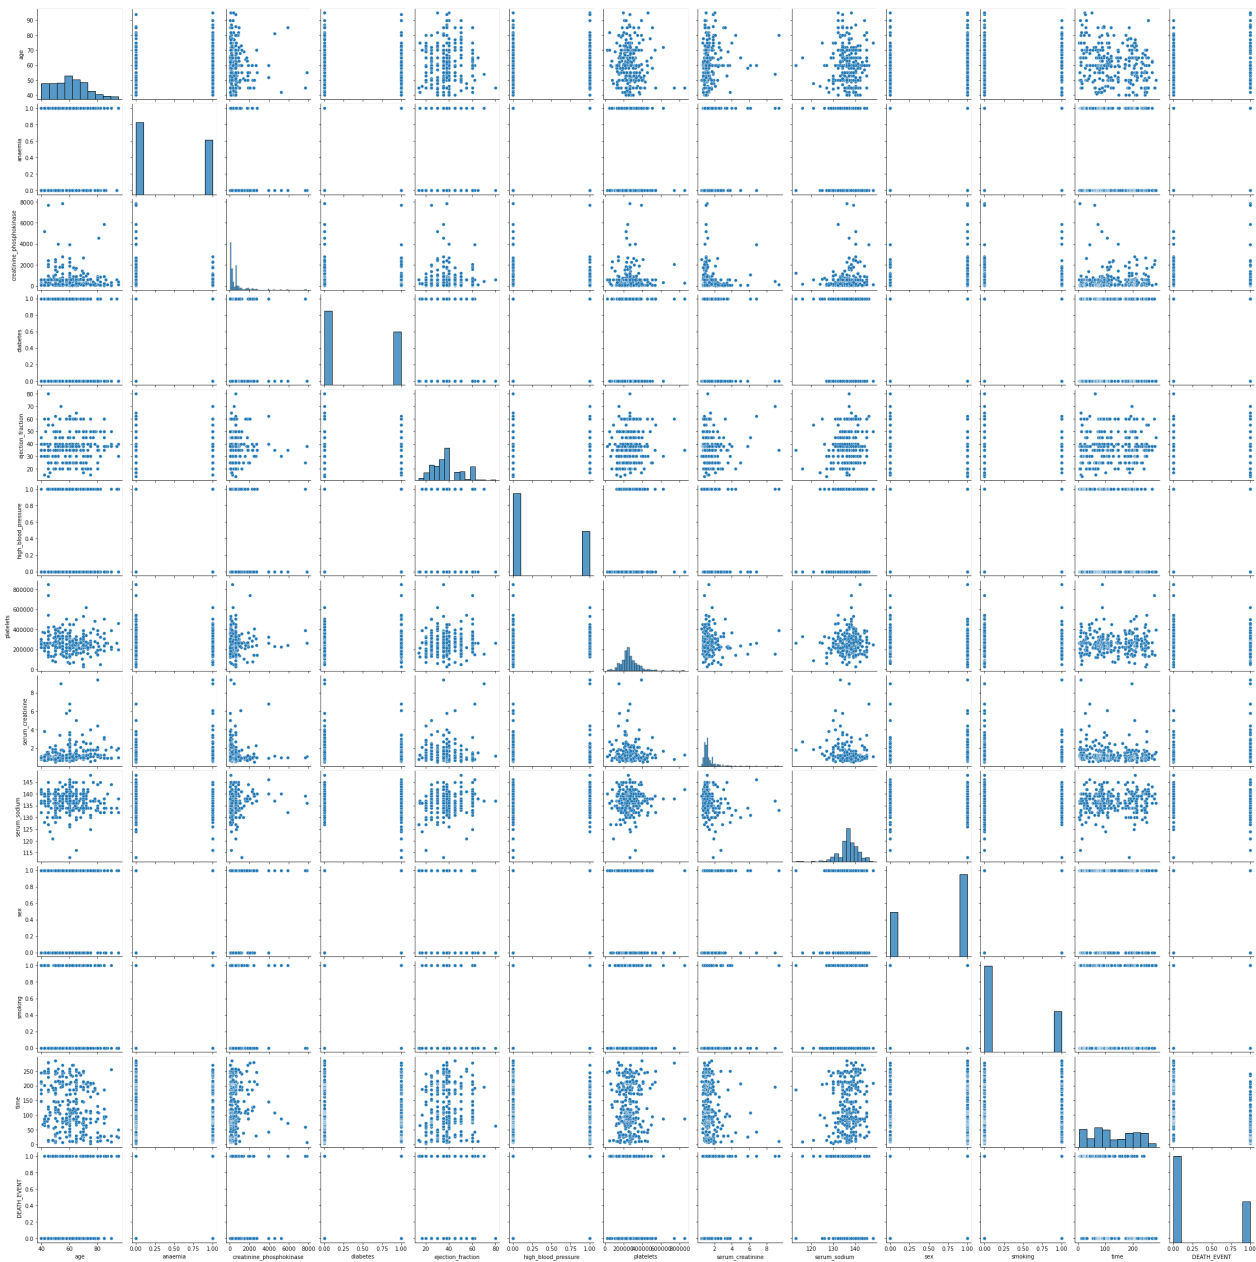
```
df.describe()
```

Out[5]:

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_pressure	platelets	serum_creatinine	sex
count	299.000000	299.000000	299.000000	299.000000	299.000000	299.000000	299.000000	299.000000	299.000000
mean	60.833893	0.431438	581.839465	0.418060	38.083612	0.351171	263358.029264	1.39388	0.488294
std	11.894809	0.496107	970.287881	0.494067	11.834841	0.478136	97804.236869	1.03451	0.501706
min	40.000000	0.000000	23.000000	0.000000	14.000000	0.000000	25100.000000	0.500000	0.000000
25%	51.000000	0.000000	116.500000	0.000000	30.000000	0.000000	212500.000000	0.900000	0.000000
50%	60.000000	0.000000	250.000000	0.000000	38.000000	0.000000	262000.000000	1.100000	0.000000
75%	70.000000	1.000000	582.000000	1.000000	45.000000	1.000000	303500.000000	1.400000	1.000000
max	95.000000	1.000000	7861.000000	1.000000	80.000000	1.000000	850000.000000	9.400000	1.000000


In [6]:

```
sns.pairplot(data=df)
plt.show()
```



In [7]:

```
plt.figure(figsize=(10,10))
sns.heatmap(df.corr(),annot=True,cmap="Blues")
plt.show()
```



In [8]:


```
df["smoking"].value_counts()
```

Out[8]:

```
0    203
1     96
Name: smoking, dtype: int64
```


In [41]:

```
plt.pie(df["smoking"].value_counts(),labels=["Those who Don't Smokes","Those who Smokes"],autopct='%1.1f%%')
plt.show()
```




In [10]:

```
sns.countplot(data=df,x="diabetes")
plt.show()
```



In [42]:

```
sns.countplot(data=df,x="high_blood_pressure")
plt.show()
```



In [12]:

```
df.drop("time",axis=1,inplace=True)
```

In [13]:

```
x=df.iloc[:,[0,3,5,10]]
y=df.iloc[:,-1]
```

In [14]:

```
x
```

Out[14]:

	age	diabetes	high_blood_pressure	smoking
0	75.0	0	1	0
1	55.0	0	0	0
2	65.0	0	0	1
3	50.0	0	0	0
4	65.0	1	0	0
...
294	62.0	1	1	1
295	55.0	0	0	0
296	45.0	1	0	0
297	45.0	0	0	1
298	50.0	0	0	1

299 rows × 4 columns

In [15]:

y

Out[15]:

```
0      1
1      1
2      1
3      1
4      1
..
294    0
295    0
296    0
297    0
298    0
Name: DEATH_EVENT, Length: 299, dtype: int64
```

In [16]:

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x=sc.fit_transform(x)
```

In [17]:

x

Out[17]:

```
array([[ 1.19294523, -0.84757938,  1.35927151, -0.68768191],
       [-0.49127928, -0.84757938, -0.73568819, -0.68768191],
       [ 0.35083298, -0.84757938, -0.73568819,  1.4541607 ],
       ...,
       [-1.33339153,  1.1798305 , -0.73568819, -0.68768191],
       [-1.33339153, -0.84757938, -0.73568819,  1.4541607 ],
       [-0.9123354 , -0.84757938, -0.73568819,  1.4541607 ]])
```

In [18]:

```
from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.3,random_state=1)
```

In [19]:

xtrain

Out[19]:

```
array([[-0.99654663, -0.84757938,  1.35927151, -0.68768191],
       [ 0.7718891 ,  1.1798305 ,  1.35927151,  1.4541607 ],
       [ 1.19294523, -0.84757938,  1.35927151, -0.68768191],
       [-0.07022315,  1.1798305 , -0.73568819, -0.68768191],
       [-0.9123354 ,  1.1798305 , -0.73568819, -0.68768191],
       [-0.74391295,  1.1798305 , -0.73568819,  1.4541607 ],
       [ 2.03505748, -0.84757938, -0.73568819, -0.68768191],
       [-1.33339153,  1.1798305 , -0.73568819, -0.68768191],
       [-0.9123354 ,  1.1798305 ,  1.35927151,  1.4541607 ],
       [-0.49127928, -0.84757938,  1.35927151, -0.68768191],
       [-1.58602521, -0.84757938, -0.73568819,  1.4541607 ],
       [-1.33339153, -0.84757938, -0.73568819,  1.4541607 ],
       [-0.65970173, -0.84757938, -0.73568819,  1.4541607 ],
       [-0.49127928, -0.84757938, -0.73568819, -0.68768191],
       [-0.15443437, -0.84757938,  1.35927151,  1.4541607 ],
       [ 0.18241053,  1.1798305 , -0.73568819, -0.68768191],
       [ 0.35083298, -0.84757938, -0.73568819, -0.68768191],
       [ 1.19294523,  1.1798305 , -0.73568819, -0.68768191]])
```

In [20]:

ytrain

Out[20]:

```

14      0
210     0
236     0
44      1
163     1
..
203     0
255     0
72      1
235     0
37      1
Name: DEATH_EVENT, Length: 209, dtype: int64

```

In [21]:

```

from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

```

In [22]:

```
from sklearn.metrics import classification_report
```

In [23]:

```

def mymodel(model):
    model.fit(xtrain,ytrain)
    ypred=model.predict(xtest)

    train=model.score(xtrain,ytrain)
    test=model.score(xtest,ytest)

    print(f"Training Acc:- {train}\n Testing Acc:- {test}")
    print(classification_report(ytest,ypred))
    return model

```

In [24]:

```
knn=mymodel(KNeighborsClassifier())
knn
```

```

Training Acc:- 0.7607655502392344
Testing Acc:- 0.5777777777777777

```

	precision	recall	f1-score	support
0	0.70	0.72	0.71	64
1	0.25	0.23	0.24	26
accuracy			0.58	90
macro avg	0.47	0.47	0.47	90
weighted avg	0.57	0.58	0.57	90

Out[24]:

KNeighborsClassifier()

In [25]:

```
mymodel(DecisionTreeClassifier())
```

```

Training Acc:- 0.8421052631578947
Testing Acc:- 0.6444444444444445

```

	precision	recall	f1-score	support
0	0.73	0.80	0.76	64
1	0.35	0.27	0.30	26
accuracy			0.64	90
macro avg	0.54	0.53	0.53	90
weighted avg	0.62	0.64	0.63	90

Out[25]:

DecisionTreeClassifier()

In [26]:

```
mymodel(RandomForestClassifier())
```

Training Acc:- 0.8421052631578947

Testing Acc:- 0.6111111111111112

	precision	recall	f1-score	support
0	0.72	0.75	0.73	64
1	0.30	0.27	0.29	26
accuracy			0.61	90
macro avg	0.51	0.51	0.51	90
weighted avg	0.60	0.61	0.60	90

Out[26]:

RandomForestClassifier()

In [27]:

```
from sklearn.model_selection import GridSearchCV
```

In [28]:

```
param={
    "n_neighbors":range(1,100),
    "leaf_size":range(1,100,10)
}
```

In [29]:

```
grid=GridSearchCV(KNeighborsClassifier(),param,verbose=2)
grid.fit(xtrain,ytrain)
```

Fitting 5 folds for each of 990 candidates, totalling 4950 fits

```
[CV] END .....leaf_size=1, n_neighbors=1; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=1; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=1; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=1; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=1; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=2; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=2; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=2; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=2; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=2; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=3; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=3; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=3; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=3; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=3; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=4; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=4; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=4; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=4; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=4; total time= 0.0s
[CV] END .....leaf_size=1, n_neighbors=4; total time= 0.0s
```

In [30]:

```
grid.best_params_
```

Out[30]:

```
{'leaf_size': 51, 'n_neighbors': 5}
```

In [31]:

```
grid.best_score_
```

Out[31]:

0.6844367015098722

In [32]:

```
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier
```

In [33]:

```
Ada=mymodel(AdaBoostClassifier())
Ada
```

```
Training Acc:- 0.7081339712918661
Testing Acc:- 0.7111111111111111
      precision    recall  f1-score   support

     0       0.73     0.94     0.82        64
     1       0.50     0.15     0.24        26

 accuracy         0.62
 macro avg         0.62
 weighted avg      0.66
```

Out[33]:

AdaBoostClassifier()

In [34]:

```
mymodel(GradientBoostingClassifier())
```

```
Training Acc:- 0.7894736842105263
Testing Acc:- 0.6555555555555556
      precision    recall  f1-score   support

     0       0.73     0.83     0.77        64
     1       0.35     0.23     0.28        26

 accuracy         0.66
 macro avg         0.54
 weighted avg      0.62
```

Out[34]:

GradientBoostingClassifier()

In [35]:

```
df.head(5)
```

Out[35]:

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_pressure	platelets	serum_creatinine	serum_sodium	sex
0	75.0	0	582	0	20	1	265000.00	1.9	130	1
1	55.0	0	7861	0	38	0	263358.03	1.1	136	1
2	65.0	0	146	0	20	0	162000.00	1.3	129	1
3	50.0	1	111	0	20	0	210000.00	1.9	137	1
4	65.0	1	160	1	20	0	327000.00	2.7	116	0

In [36]:

```
def predictDeathEvent():
    age=float(input("Enter the Age of the person:- "))
    diabetes=int(input("Enter If you have Diabetes or not:- "))
    high_blood_pressure=int(input("Enter your If you have High Blood Pressure:- "))
    smoking=int(input("Enter If you Smoke or not:- "))

    newob=[[age,diabetes,high_blood_pressure,smoking]]
    ypred=Ada.predict(newob)
    return ypred[0]
```

In [37]:

```
predictDeathEvent()
```

```
Enter the Age of the person:- 20
Enter If you have Diabetes or not:- 1
Enter your If you have High Blood Pressure:- 1
Enter If you Smoke or not:- 1
```

Out[37]:

1

In []: